

LECTURE 3: Optimization in Higher Dimensions

- Optimization (maximization/minimization) is of huge importance in data analysis and is the basis for recent breakthroughs in machine learning and big data
- A lot of it is application dependent and there is a vast number of methods developed: we cannot cover them all in this lecture

LECTURE 3: Optimization in Higher Dimensions


- Broadly can be divided into 1st order (derivatives are available, but not Hessian) and 2nd order (approximate Hessian or full Hessian evaluation)
- 0th order: no gradients available: use finite difference to get the gradient or use downhill simplex (Nelder & Mead method). Very slow and we will not discuss them here.

Preparation of Parameters

- Often the parameters are not unconstrained: they may be positive (or negative), or bounded to an interval
- First step is to make optimization unconstrained: map the parameter to a new parameter that is unbounded. For example, if a variable is positive, $x > 0$, use $z = \log(x)$ instead of x .
- One can also change the prior so that it reflects the original prior: $p_{pr}(z)dz = p_{pr}(x)dx$
- If $x > 0$ has uniform prior in x then $p_{pr}(z) = dx/dz = x = e^z$

General Strategy

- We want to descend down a function $J(a)$ (if minimizing) using iterative sequence of steps at a_t . For this we need to choose a direction p_t and move in that direction: $J(a_t + \eta p_t)$
- A few options: fix η
- line search: vary η until $J(a_t + \eta p_t)$ is minimized
- Trust region: construct an approximate quadratic model for J and minimize it but only within trust region where quadratic model is approximately valid



Line Search Directions and Backtracking

- Gradient descent: Gradient - $\nabla_a J(a, x_t)$
- Newton: Inverse Hessian H^{-1} times gradient
 $-H^{-1} \nabla_a J(a)$
- Quasi-Newton: approximate H^{-1} with B^{-1} (SR1 and BFGS)
- Nonlinear conjugate gradient:
 $p_t = -\nabla_a J(a, x_t) + \beta_t p_{t-1}$, where p_{t-1} and p_t are conjugate
- Step length with backtracking: choose first proposed length
- If it does not reduce the function value reduce it by some factor, check again
- Repeat until step length is ϵ , at that point switch to gradient descent

Trust Region Method

- **Multi-dim parabola method**: define approximate quadratic function, but limit the step


$$\min_{p \in \mathbb{R}^n} m_k(p) = f_k + g_k^T p + \frac{1}{2} p^T B_k p \quad \text{s.t. } \|p\| \leq \Delta_k$$

- Here Δ_k is **trust region radius**
- Evaluate at previous iteration and compare the actual reduction to predicted reduction

$$\rho_k = \frac{f(x_k) - f(x_k + p_k)}{m_k(0) - m_k(p_k)}$$

- If ρ_k around 1 we can increase Δ_k
- If close to 0 or negative we shrink Δ_k


- If trust region covers m center step there
- Otherwise direction of step changes



Constrained Optimization: Lagrange Multiplier Method

- If the center of m is inside trust region step there
- Otherwise we must solve constrained optimization
- We solve this optimization with Lagrange multiplier method:
minimize $f + g^T p + p^T B p + \lambda (p^2 - \Delta^2)$ with respect to p and λ .
Gradient w.r.t. λ gives the constraint $p^2 = \Delta^2$, thus the constraint is automatically satisfied. This determines the value of λ .
- Minimization with respect to p now includes λp^2
- As a result the step direction is not towards center of m when trust region does not cover it: see picture on next slide

Line Search vs. Trust Region



1st Order: Gradient Descent


- We have a vector of parameters \mathbf{a} and a scalar loss (cost) function $J(\mathbf{a}, x, y)$ which is a function of a data vector (x, y) we want to optimize (say minimize). This could be a nonlinear least square loss function: $J = \chi^2$

$$\chi^2(\mathbf{a}) = \sum_{i=0}^{N-1} \left[\frac{y_i - y(x_i | \mathbf{a})}{\sigma_i} \right]^2$$

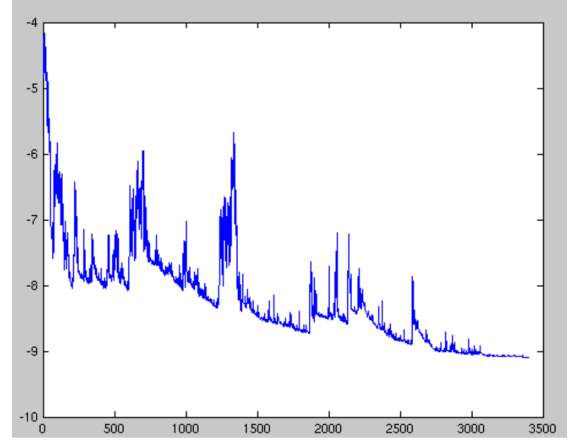
- (Batch) gradient descent updates all the variables at once:
 $\delta \mathbf{a} = - \eta \nabla_{\mathbf{a}} J(\mathbf{a})$: in ML. η is called learning rate
- It gets stuck on saddle points, where gradient is 0 everywhere (see animation later)

Scaling

- Change variables to make surface more circular
- Example: change of dimensions



Stochastic Gradient Descent


- Stochastic gradient descent: do this just for one data pair x_i, y_i :
$$\delta a = -\eta \nabla_a J(a, x_i, y_i)$$
 - This saves on computational cost, but is noisy, so one repeats it by randomly choosing data i
 - Has large fluctuations in the cost function
- 
- This is potentially a good thing: it may avoid getting stuck in the local minima (or saddle points)
 - Learning rate is slowly reduced
 - Has revolutionized machine learning

Mini-batch Stochastic Gradient

- Mini-batch takes advantage of hardware and software implementations where a gradient w.r.t. to a number of data points can be evaluated as fast as a single data (e.g. mini-batch of $N = 256$)
- Challenges of (stochastic) gradient descent: how to choose learning rate (in 2nd order methods this is given by Hessian)
- Ravines:



Ravines




(a)



(b)


Adding Momentum: Rolling down the hill

- We can add momentum and mimic a ball rolling down the hill
- Use previous update as the direction
- $v_t = \gamma v_{t-1} + \eta \nabla_a J(a)$, $da = -v_t$ with γ of order 1 (e.g. 0.9)
- Momentum increases for directions where gradient does not change



Nesterov Accelerated Gradient


- We can predict where to evaluate next gradient using previous velocity update
- $v_t = \gamma v_{t-1} + \eta \nabla_a J(a - \gamma v_{t-1}), \delta a = -v_t$
- Momentum (blue) vs NAG (brown+red=green)



- See <https://arxiv.org/abs/1603.04245> for theoretical justification of NAG based on a Bregman divergence Lagrangian

Adagrad, Adadelata, Rmsprop, ADAM, ...

- Make the learning rate h dependent on a_i
- Use past gradient information to update h
- Example ADAM: ADAptive Momentum estimation
- $m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$ $g_t = \nabla_a J(a)$
- $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$
- bias correction: $m_t' = m_t / (1 - \beta_1)$, $v_t' = v_t / (1 - \beta_2)$
- Update rule: $\delta a = -\eta / (v_t'^{1/2} + \varepsilon)$
- Recommended values $\beta_1 = 0.9$, $\beta_2 = 0.999$, $\varepsilon = 10^{-8}$
- The methods are empirical (show animation)



2nd Order Method: Newton

- We have seen that there is no natural way to choose learning rate in 1st order methods
- But Newton's method provides a clear answer what the learning rate should be:
- $J(a+\delta a) = J(a) + \delta a \nabla_a J(a) + \delta a \delta a' \nabla_a \nabla_a J(a)/2 \dots$
- Hessian $H_{ij} = \nabla_{a_i} \nabla_{a_j} J(a)$
- At the extremum we want $\nabla_a J(a) = 0$ so a Newton update step is $\delta a = -H^{-1} \nabla_a J(a)$
- We do not need to guess the learning rate
- We do need to evaluate Hessian and invert it (or use LU): expensive in many dimensions!
- In many dimensions we use iterative schemes to solve this problem

Quasi-Newton

- Computing Hessian and inverting it is expensive, but one can approximate it with a low rank tensor
- Symmetric rank 1 (SR1) $s_k = x_{k+1} - x_k, \quad y_k = \nabla f_{k+1} - \nabla f_k$

$$B_{k+1} = B_k + \frac{(y_k - B_k s_k)(y_k - B_k s_k)^T}{(y_k - B_k s_k)^T s_k} \quad B_{k+1} s_k = y_k$$

- BFGS (rank 2 update, positive definite)

$$B_{k+1} = B_k - \frac{B_k s_k s_k^T B_k}{s_k^T B_k s_k} + \frac{y_k y_k^T}{y_k^T s_k}$$

- Inverse (Woodbury formula)

$$B_{k+1}^{-1} = (I - \rho_k s_k y_k^T) B_k^{-1} (I - \rho_k y_k s_k^T) + \rho_k s_k s_k^T, \quad \rho_k = \frac{1}{y_k^T s_k}$$

L-BFGS



- For large problems this gets too expensive. Limited memory BFGS updates only based on last N iterations (N of order 10-100)
- In practice increasing N often does not improve the results
- Historical note: quasi-Newton methods originate from W.C. Davidon's work in 1950s, a physicist at Argonne national lab.

Linear Conjugate Direction

- Is an iterative method to solve $\mathbf{Ax} = \mathbf{b}$ (so belongs to linear algebra)
- Can be used for optimization: $\min \mathbf{J} = \mathbf{x}^T \mathbf{Ax} - \mathbf{b}^T \mathbf{x}$
- Conjugate vectors: $\mathbf{p}_i^T \mathbf{A} \mathbf{p}_j = 0$ for all i, j not equal i
- Construction similar to Gram-Schmidt (QR), where \mathbf{A} plays the role of scalar product norm: $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$ where $\alpha_k = -\mathbf{r}_k^T \mathbf{p}_k / (\mathbf{p}_k^T \mathbf{A} \mathbf{p}_k)$ and $\mathbf{r}_k = \mathbf{Ax}_k - \mathbf{b}$
- Essentially we are taking a dot product (with \mathbf{A} norm) of the vector with previous vectors to project it perpendicular to previous vectors
- Since the space is N -dim after N steps we have spanned the full space and converged to true solution, $\mathbf{r}_N = 0$.


Conjugate Direction

- If we have the matrix \mathbf{A} in diagonal form so that basis vectors are orthogonal we can find the minimum trivially along the axes, otherwise not



Linear Conjugate Gradient

- Computes \mathbf{p}_k from \mathbf{p}_{k-1}
- We want the step to be linear combination of residual $-\mathbf{r}_k$ and previous direction \mathbf{p}_{k-1}
- $\mathbf{p}_k = -\mathbf{r}_k + \beta_k \mathbf{p}_{k-1}$ premultiply by $\mathbf{p}_{k-1}^T \mathbf{A}$
- $\beta_k = (\mathbf{r}_k^T \mathbf{A} \mathbf{p}_{k-1}) / (\mathbf{p}_{k-1}^T \mathbf{A} \mathbf{p}_{k-1})$ imposing $\mathbf{p}_{k-1}^T \mathbf{A} \mathbf{p}_k = 0$
- Converges rapidly for similar eigenvalues, not so much if condition number is high



Preconditioning

- Tries to improve condition number of \mathbf{A} by multiplying by another matrix \mathbf{C} that is simple

$$\hat{x} = Cx.$$

$$\hat{\phi}(\hat{x}) = \frac{1}{2}\hat{x}^T (C^{-T} A C^{-1}) \hat{x} - (C^{-T} b)^T \hat{x}.$$

$$(C^{-T} A C^{-1}) \hat{x} = C^{-T} b$$


- We wish to reduce condition number of $C^{-T} A C^{-1}$
- Example: incomplete Cholesky $\mathbf{A} = \mathbf{L}\mathbf{L}^T$ by computing only a sparse \mathbf{L}
- Preconditioners are very problem specific

Nonlinear Conjugate Gradient

- Replace \mathbf{a}_k with line search that minimizes \mathbf{J} , and use $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{a}_k \mathbf{p}_k$
- Replace $\mathbf{r}_k = \mathbf{A}\mathbf{x}_k - \mathbf{b}$ with gradient of \mathbf{J} : $\nabla_a \mathbf{J}$
- This is Fletcher-Reeves version, Polak-Ribiere modifies β
- CG is one of the most competitive methods, but requires the Hessian to have low condition number
- Typically we do a few CG steps at each k , then move on to a new gradient evaluation

CG vs. Gradient Descent

- In 2-d CG has to converge in 2 steps



Gauss-Newton for Nonlinear Least Squares

$$\chi^2(\mathbf{a}) = \sum_{i=0}^{N-1} \left[\frac{y_i - y(x_i|\mathbf{a})}{\sigma_i} \right]^2$$

$$\frac{\partial \chi^2}{\partial a_k} = -2 \sum_{i=0}^{N-1} \frac{[y_i - y(x_i|\mathbf{a})]}{\sigma_i^2} \frac{\partial y(x_i|\mathbf{a})}{\partial a_k} \quad k = 0, 1, \dots, M-1$$

$$\frac{\partial^2 \chi^2}{\partial a_k \partial a_l} = 2 \sum_{i=0}^{N-1} \frac{1}{\sigma_i^2} \left[\frac{\partial y(x_i|\mathbf{a})}{\partial a_k} \frac{\partial y(x_i|\mathbf{a})}{\partial a_l} - [y_i - y(x_i|\mathbf{a})] \frac{\partial^2 y(x_i|\mathbf{a})}{\partial a_l \partial a_k} \right]$$

$$\beta_k \equiv -\frac{1}{2} \frac{\partial \chi^2}{\partial a_k} \quad \alpha_{kl} \equiv \frac{1}{2} \frac{\partial^2 \chi^2}{\partial a_k \partial a_l} \quad \sum_{l=0}^{M-1} \alpha_{kl} \delta a_l = \beta_k$$

$$\alpha_{kl} = \sum_{i=0}^{N-1} \frac{1}{\sigma_i^2} \left[\frac{\partial y(x_i|\mathbf{a})}{\partial a_k} \frac{\partial y(x_i|\mathbf{a})}{\partial a_l} \right]$$

We drop 2nd term in Hessian because residual $r = y_i - y$ is small, fluctuates around 0 and because y'' may be small (or zero for linear problems)

Line search in direction $\delta \mathbf{a}$

Gauss-Newton + Trust Region = Levenberg-Marquardt Method

- Solving $\mathbf{A}^T \mathbf{A} \delta \mathbf{a} = \mathbf{A}^T \mathbf{b}$ is equivalent to minimize $|\mathbf{A} \delta \mathbf{a} - \mathbf{b}|^2$
- if trust region is within the solution just solve this equation
- If not we need to impose $\|\delta \mathbf{a}\| = \Delta_k$
- Lagrange multiplier minimization equivalent to $(\mathbf{A}^T \mathbf{A} + \lambda \mathbf{I}) \delta \mathbf{a} = \mathbf{A}^T \mathbf{b}$ and $\lambda(\Delta - \|\delta \mathbf{a}\|) = 0$
- For small λ this is Gauss-Newton (use close to minimum), for large λ this is steepest descent (use far from minimum)
- A good method for nonlinear least squares

Summary

- Optimization one of key numerical methods of modern data analysis. Typical examples are nonlinear least square problem and ML parameters (e.g. neural networks etc.)
- If at this point you are confused which methods you should use you are not alone: it depends on application and often the best way to answer is to try
- Some general guidances: if there is a lot of parameters (e.g. ML) and likelihood evaluations are cheap then use 1st order methods

Summary

- If the data is independent and there is a lot of data then use stochastic 1st order methods, e.g. ADAM
- If the likelihood evaluation is slow and number of parameters low use Newton or Gauss-Newton (e.g. Levenberg-Marquardt)
- If likelihood slow and number of parameters large use approximate Newton or Gauss-Newton (e.g. Steihaug with nonlinear CG), or use quasi-Newton (e.g. L-BFGS)
- Choosing a method is not enough: you also need to choose line search method (e.g. backtracking, Wolfe conditions) or trust region determination
- Typically these methods only find local minimum. Non-convex problems are hard: we will look at some stochastic methods (e.g. simulated annealing) in next lecture

Literature

- *Numerical Recipes*, Press et al., Chapter 9, 10, 15
- *Computational Physics*, M. Newman, Chapter 6
- Nocedal and Wright, Optimization
- <https://arxiv.org/abs/1609.04747>